

The AI Capital Supercycle

OpenAI's move toward a trillion-dollar IPO turns the largest infrastructure bet in tech history into a public-markets referendum — just as falling model prices begin to question whether the spending will ever pay for itself.

MID-TERM

EXECUTIVE BRIEF

- OpenAI is preparing to **confidentially file** for a public listing targeted **above \$1 trillion**, with Goldman Sachs and Morgan Stanley advising and a Q4 2026 listing in view.
- It enters the process off a record **\$122 billion raise** at a roughly **\$852 billion** post-money valuation, against about **\$25 billion** in annualized revenue in early 2026.
- On the other side of the ledger sit more than **\$1.4 trillion** in committed data-center and infrastructure spending, an estimated **\$27 billion** of 2026 cash burn, and no expected profit before roughly **2030**.
- The US macro backdrop is supportive: the **Dow closed at a record 50,285.66** and the Nasdaq logged its strongest week since November, led by AI megacaps.
- The bear case sharpened this week — analysts warn that **rapidly cheapening models**, partly from efficient open-weight rivals, could compress the margins the valuations assume.
- Adjacent moves confirm the build-out: **Anthropic is in early talks** to use Microsoft's custom **Maya 200** chips, while Microsoft Agent 365 and Google's 24/7 Gemini Spark agent reached the market.

THE SIGNAL LANDSCAPE

The Filing Heard Around the Market

In late May, CNBC and The Wall Street Journal reported that OpenAI is preparing to confidentially submit a draft IPO prospectus, with Goldman Sachs and Morgan Stanley steering the process toward a listing valued above \$1 trillion. That would rank among the largest technology debuts ever and, more importantly, would be the first time the company most synonymous with the generative-AI boom is priced by public investors rather than a small circle of private backers.

The filing follows a record \$122 billion funding round that set the post-money valuation near \$852 billion. Crossing the trillion-dollar line on the public tape would not merely reward existing shareholders; it would set a market-wide benchmark for how every other AI company — from Anthropic to a queue of infrastructure providers — is valued. The IPO is therefore less a single event than a stress test for the entire asset class.

Source: CNBC / The Wall Street Journal, May 2026

The Arithmetic of \$1.4 Trillion

The case rests on a vertiginous gap between today's revenue and tomorrow's bills. OpenAI exited 2025 above \$20 billion in annualized revenue and reached roughly \$25 billion by February 2026 — extraordinary growth by any normal standard, yet small next to the more than \$1.4 trillion the company has committed to data centers and compute over the coming years. Internal projections do not foresee profitability before the end of the decade, with 2026 cash burn estimated near \$27 billion.

Public investors will be asked to underwrite a decade-long build-out on the promise that AI becomes utility-grade infrastructure embedded in every workflow. That is a coherent thesis, but it converts a research lab into a capital-intensive utility whose returns depend on durable pricing power. The question the IPO forces into the open is whether such pricing power exists in a market racing toward commoditization.

Source: OpenAI disclosures; Sacra; Investing.com, 2026

The Cheap-AI Counterweight

This week supplied the clearest counter-narrative yet. CNBC reported that rapidly falling model prices — driven partly by highly efficient open-weight competitors and an industry-wide pivot from raw scale to tokens-per-dollar — could derail the economics underpinning the OpenAI and Anthropic listings. May's notable releases reinforced the point: Google's Gemini 3.5 Flash reached general availability at about four times the speed of comparable models, and subquadratic architectures pushed context windows toward the multi-million-token range at lower cost.

If frontier capability keeps commoditizing, the enormous infrastructure outlays may outrun the revenue needed to service them — the structural risk behind louder 'AI bubble' warnings. The bull and bear cases now hinge on the same variable: whether efficiency gains expand the market faster than they erode per-token margins.

Source: CNBC; WhatLLM, May 2026

From Models to Agents (and Chips)

The supercycle is broadening beyond chatbots. Microsoft's Agent 365 — a control plane for fleets of workplace agents — became generally available, and Google unveiled an always-on, 24/7 Gemini Spark agent, signaling that 'an agent per worker' is becoming the assumed unit of enterprise software. Demand for compute is, in turn, reshaping silicon: Anthropic is in early talks to deploy Microsoft's custom Maya 200 accelerators, part of a wider scramble to reduce dependence on a single chip supplier.

Together these moves explain why the spending is so aggressive. The prize is not a better answer engine but a layer of autonomous software that runs continuously and acts on a user's behalf — a far larger market, and a far larger compute bill, than conversational AI alone.

Source: Tech Startups; CNBC, May 2026

WHAT TO WATCH

- Whether OpenAI's confidential draft converts into a public S-1 with a firm Q4 2026 listing date and a final valuation band.
- Token prices and gross margins across frontier providers — the single cleanest read on whether 'cheap AI' is eroding the thesis.
- Custom-silicon deals (Anthropic–Microsoft Maya 200 and others) that could loosen Nvidia's pricing grip and reset infrastructure economics.
- Enterprise agent-adoption metrics now that Microsoft Agent 365 and Gemini Spark are generally available — usage, not announcements.
- Market breadth: whether the record-setting rally widens beyond a handful of AI megacaps or stays dangerously concentrated.

THE OPPORTUNITY

The most resilient positions sit in the infrastructure the supercycle must buy no matter which model wins — power generation and grid interconnects, data-center cooling, custom inference silicon, and the security, identity, and orchestration layers that make autonomous agents safe to deploy. These 'picks and shovels' get paid whether AI margins expand or compress, and several are supply-constrained enough to hold pricing power through a valuation reset. For operators, the parallel move is to pilot agentic workflows now while model prices fall — locking in productivity gains before competitors do, without betting the business on any single vendor surviving the shakeout.